




TOON paper-review validation

Same 60-sample pass rate, 68.2% fewer total tokens

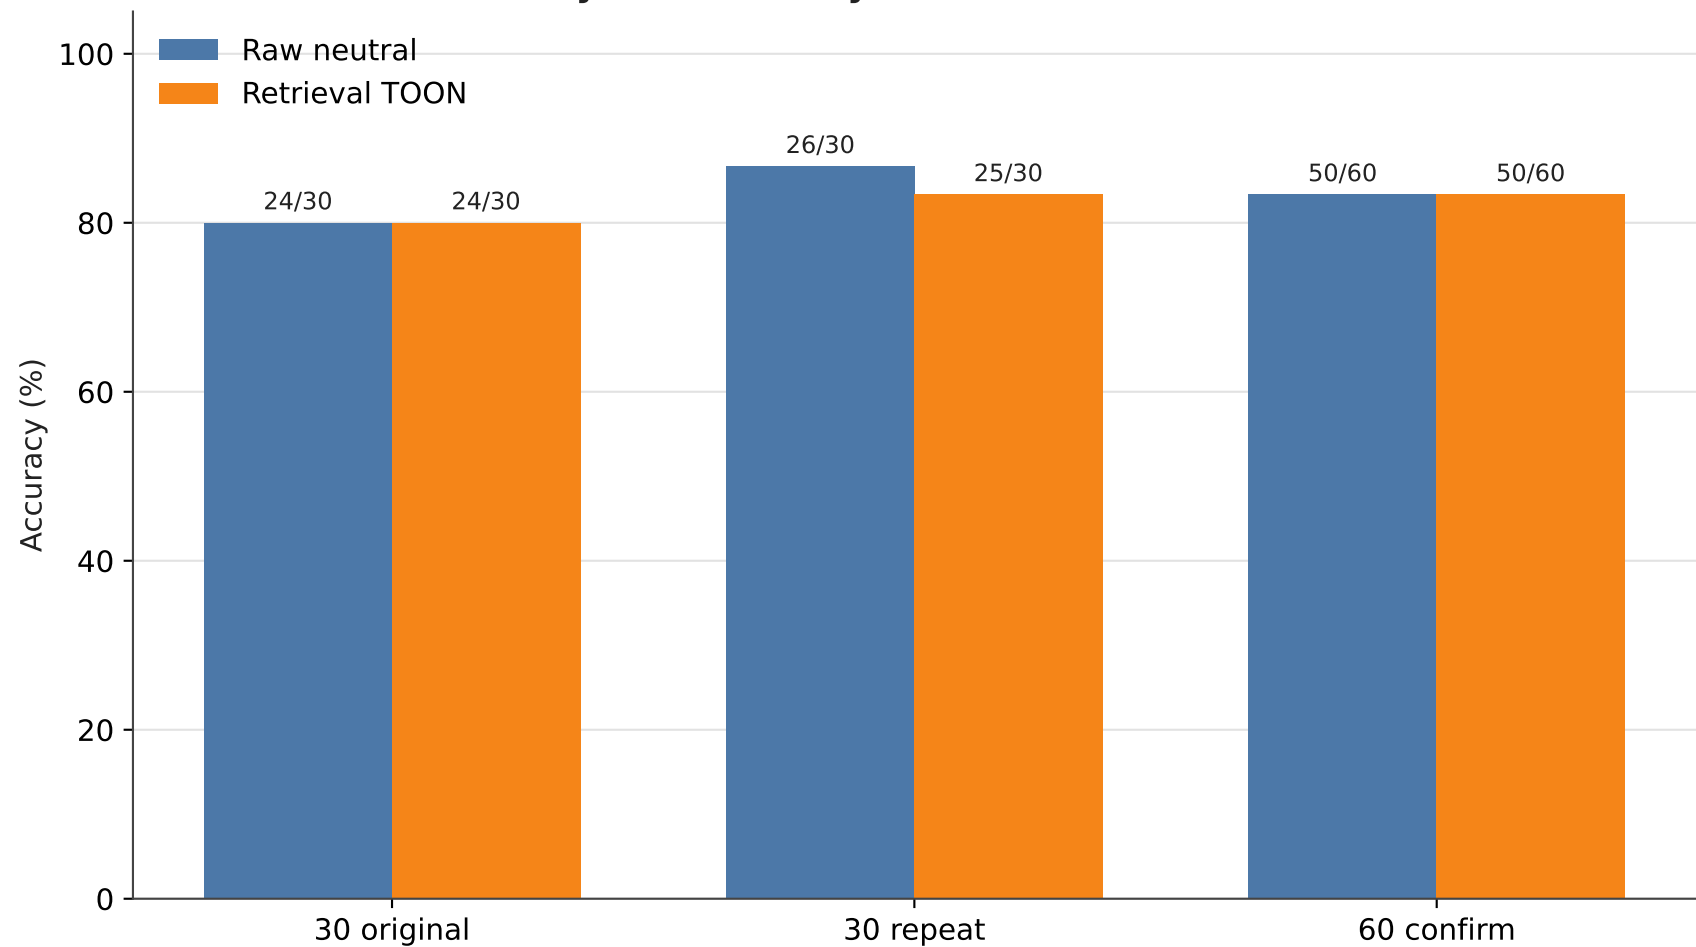


Controlled comparisons

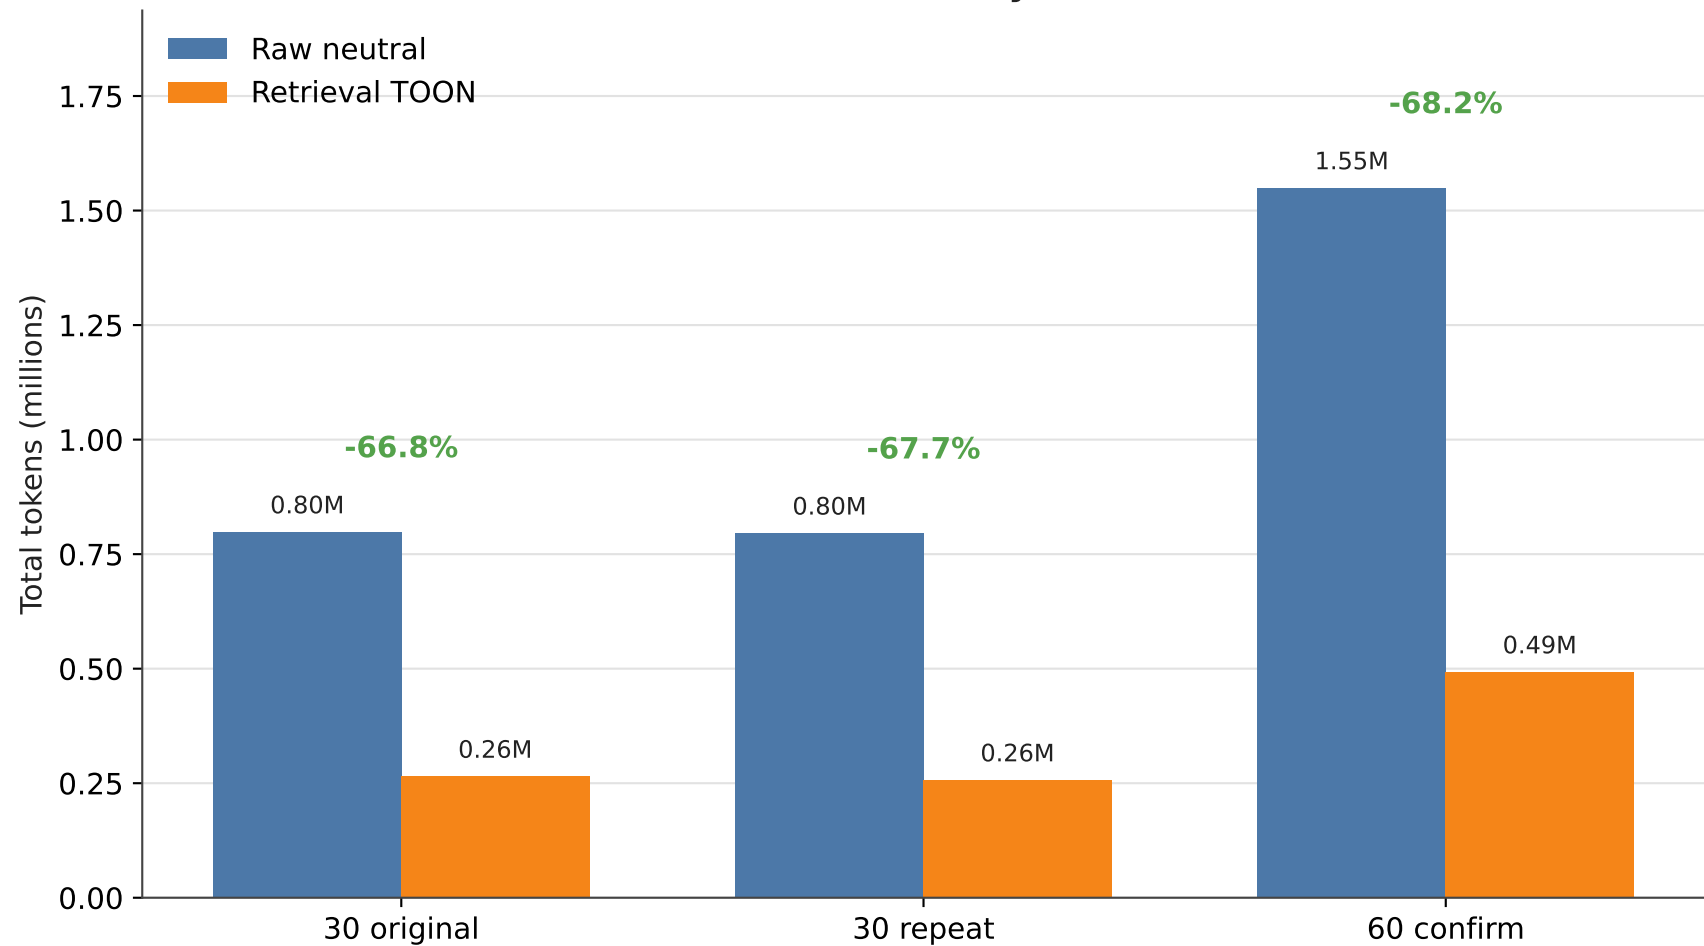
Run	Raw	TOON	Accuracy delta	Total tokens
30-sample original	24/30	24/30	+0.0000	-66.84%
30-sample repeat	26/30	25/30	-0.0333	-67.71%
60-sample confirmation	50/60	50/60	+0.0000	-68.21%

Bottom line: no aggregate accuracy damage was observed; residual risk is distributional, with TOON weaker on short papers in the 60-sample slice.

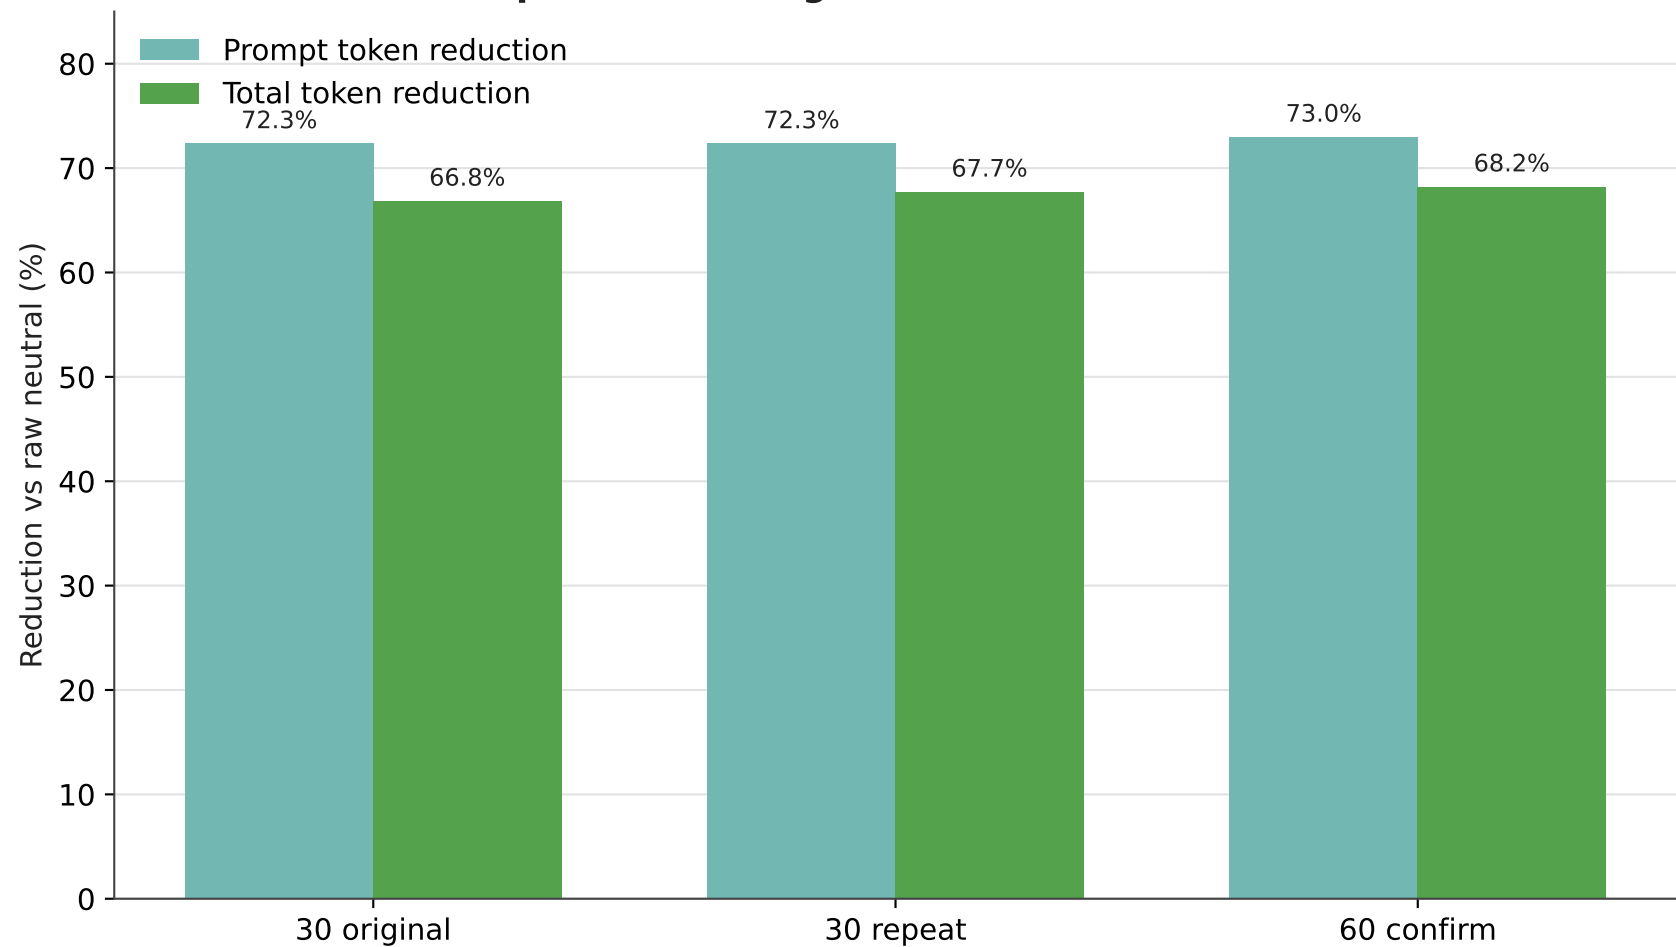
Pass rate stayed effectively matched across validation runs



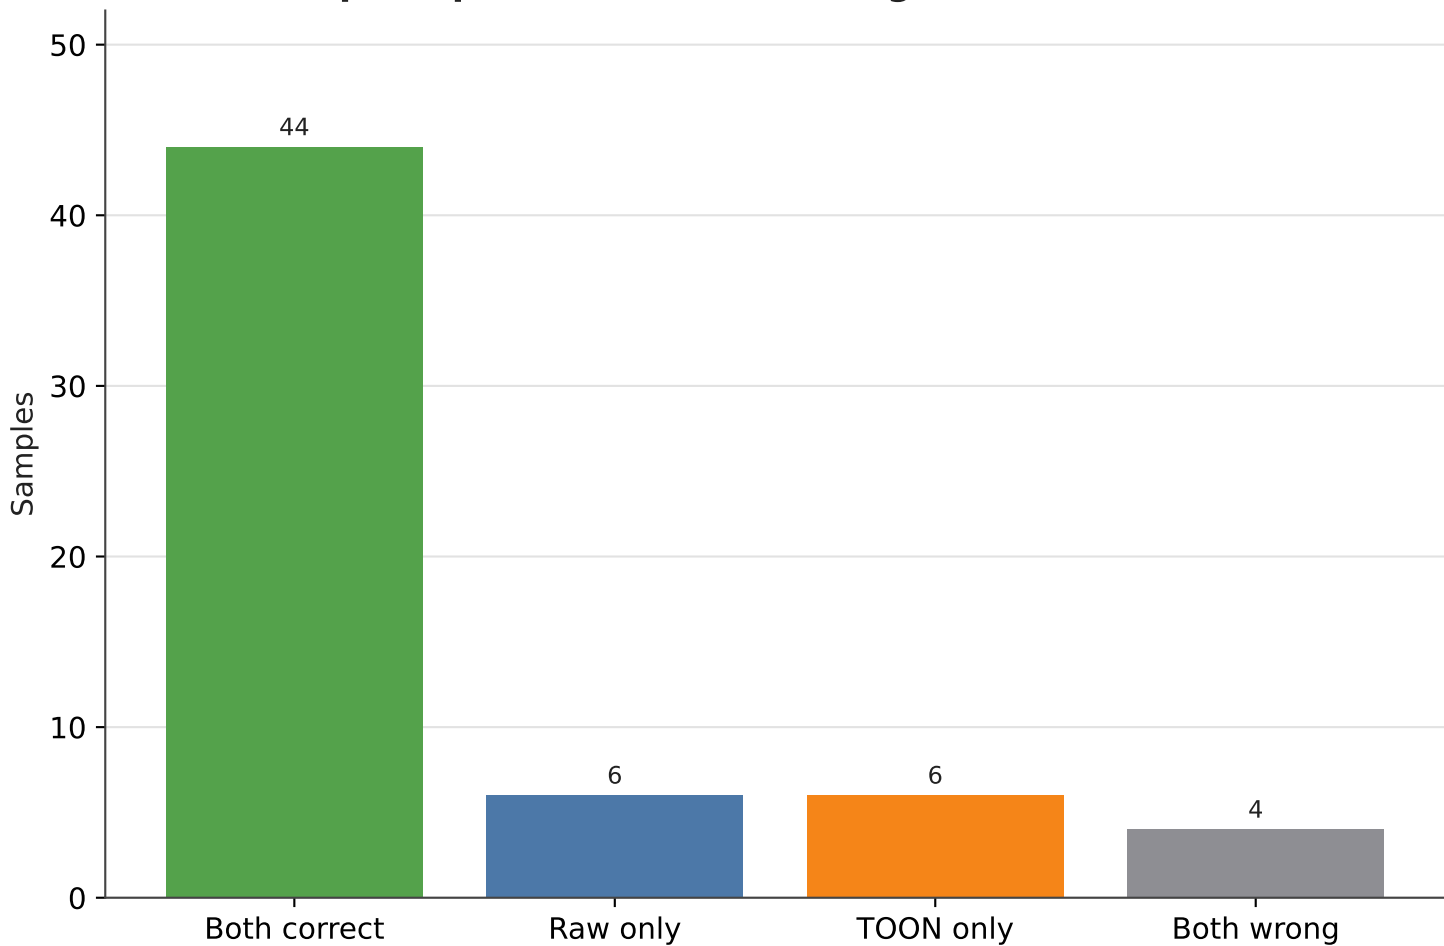
TOON cut total token use by about two thirds



Prompt-token savings were stable across reruns

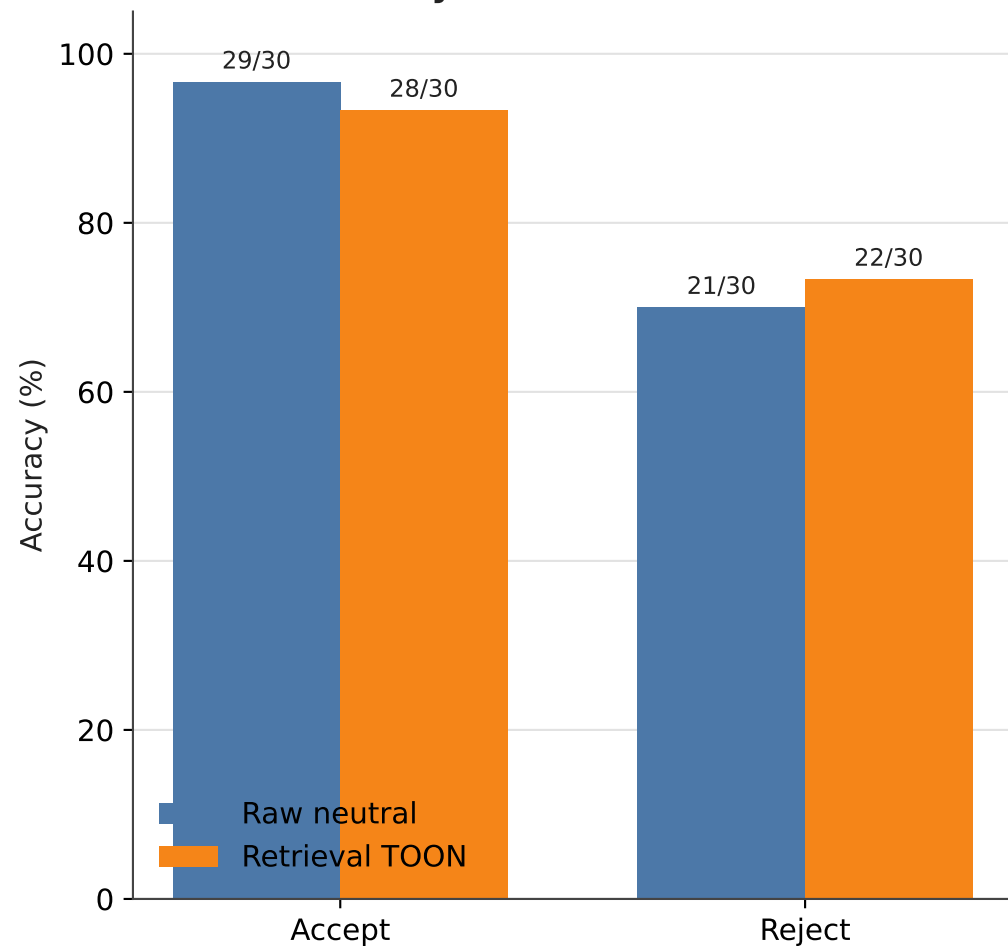


60-sample flips were balanced: 6 against TOON, 6 in favor



60-sample stratified accuracy

By decision label



By paper length

